

## Chapter 5 –Regression

Instructor: Dr. Hongfu Liu  
Email: hongfuliu@brandeis.edu

Answer “What is the relationship between the variables?”

- Dependent (response) variable – what to be predicted.
- One or more numerical or categorical independent (explanatory) variables.

Given a dataset  $\{\langle \vec{x}_n, y_n \rangle\}_{n=1, \dots, N}$ , where  $\vec{x}_n \in \mathcal{R}^d$  and  $y_n \in \mathcal{R}$

Build a function  $y = f(\vec{x}, \vec{w})$  to predict the output  $y$  of a given sample  $\vec{x}$ , where  $\vec{w}$  is the parameter of the function  $f(\cdot)$

# Simple Linear Regression



## Systolic Blood Pressure (SBP)

Age	SBP (mm Hg)
22	131
23	128
24	116
27	106
28	114
29	123
30	117
32	122
33	99
35	121
40	147

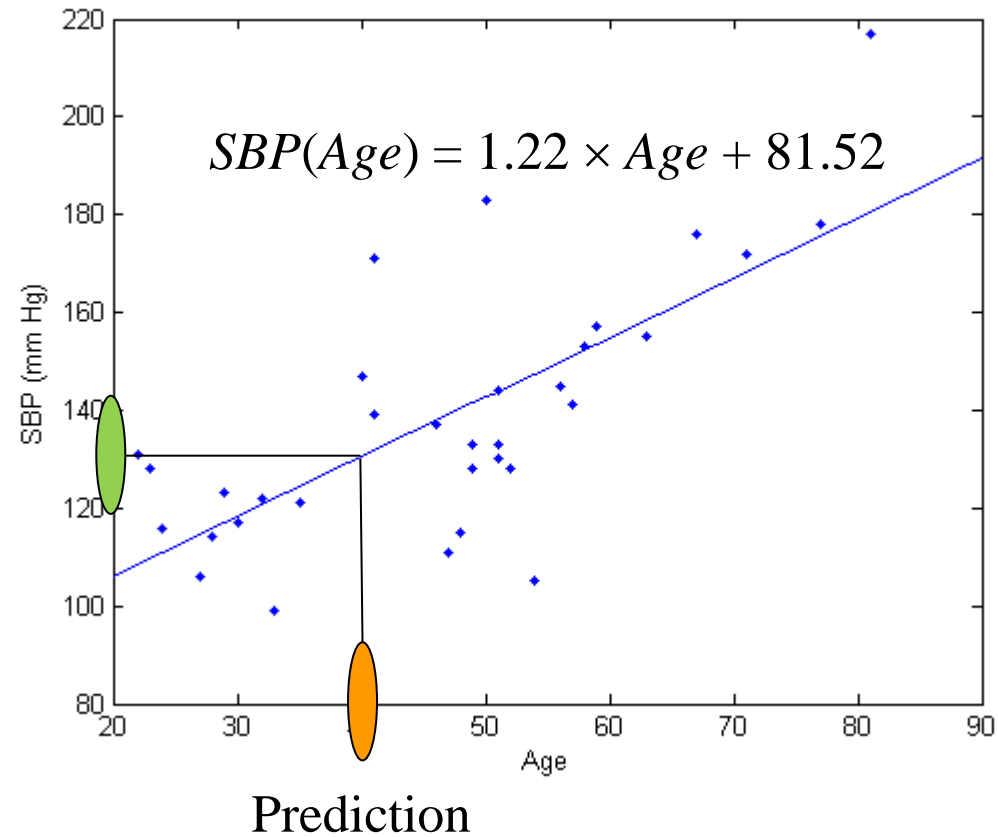
Age	SBP (mm Hg)
41	139
41	171
46	137
47	111
48	115
49	133
49	128
50	183
51	130
51	133
51	144

Age	SBP (mm Hg)
52	128
54	105
56	145
57	141
58	153
59	157
63	155
67	176
71	172
77	178
81	217

# Simple Linear Regression



$f(x) = ax + b$ , where  $a$   
measures the association  
between  $y$  and  $x$



A linear combination of the input variables

$$f(\vec{x}, \vec{w}) = w_0 + w_1x_1 + \cdots + w_dx_d$$

$$\text{Let } \vec{w} = \begin{bmatrix} w_0 \\ w_1 \\ \vdots \\ w_d \end{bmatrix} \quad \vec{x} = \begin{bmatrix} 1 \\ x_1 \\ \vdots \\ x_d \end{bmatrix}$$

$$f(\vec{x}, \vec{w}) = \vec{w}^T \vec{x}$$

Given  $\{\langle \vec{x}_n, y_n \rangle\}_{n=1, \dots, N}$  and  $f(\vec{x}, \vec{w}) = \vec{w}^T \vec{x}$

Define the objective  $Err = \sum_{n=1}^N (y_n - f(\vec{x}_n, \vec{w}))^2$

Arrange the data into the following

$$\vec{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{bmatrix} \quad X = \begin{bmatrix} \vec{x}_1^T \\ \vec{x}_2^T \\ \vdots \\ \vec{x}_N^T \end{bmatrix} = \begin{bmatrix} 1 & x_{11} & x_{12} & \cdots & x_{1d} \\ 1 & x_{21} & x_{22} & \cdots & x_{2d} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & x_{N1} & x_{N2} & \cdots & x_{Nd} \end{bmatrix}$$

Obtain  $Err = (\vec{y} - X\vec{w})^T (\vec{y} - X\vec{w})$

Solution (if  $X$  is full rank):  $\vec{w} = (X^T X)^{-1} X^T \vec{y}$       How?

Minimize  $Err = (\vec{y} - X\vec{w})^T (\vec{y} - X\vec{w})$

$$\frac{\partial Err}{\partial \vec{w}} = -2X^T (\vec{y} - X\vec{w})$$

$$\frac{\partial^2 Err}{\partial \vec{w}^2} = 2X^T X$$

Assuming  $X$  is full rank

$$\frac{\partial Err}{\partial \vec{w}} = 0 \quad \Rightarrow \quad \vec{w} = (X^T X)^{-1} X^T \vec{y}$$

$$\vec{w} = (X^T X)^{-1} X^T \vec{y} \quad \text{What if } (X^T X) \text{ is singular?}$$

How to deal with categorical variables?

Create dummy variables, one for each categorical value, and delete one dummy variable.



## Coefficient of Determination $R^2$

$R^2$  is the proportion of the variation in the dependent variable that is predictable from the independent variable(s).

$$R^2 = 1 - \frac{SS_{residual}}{SS_{total}}$$

$$SS_{residual} = Err = \sum_{n=1}^N (y_n - f(\vec{x}_n, \vec{w}))^2$$

$$SS_{total} = \sum_{n=1}^N (y_n - \bar{y})^2 \quad \text{where } \bar{y} = \frac{1}{N} \sum_{n=1}^N y_n$$

equivalent to

$$f(\vec{x}, \vec{w}) = \bar{y} + 0 \cdot x_1 + \cdots + 0 \cdot x_d$$

$$R^2 = 1 - \frac{SS_{residual}}{SS_{total}}$$

$R^2$ : The larger or smaller, the better?

- Very often, the input variables are of very different scale.
- Normalize all input variables so they are “comparable” in scale.
  - One popular way: apply a linear transformation to each variable
- Comparable scales lead to more stable numerical performance
- Comparable scales make it easier to interpret coefficients

# Linear Basis Function Model

$$f(\vec{x}, \vec{w}) = w_0 + w_1 x_1 + \cdots + w_d x_d$$

linear  
relationship  
input  $\longleftrightarrow$  output

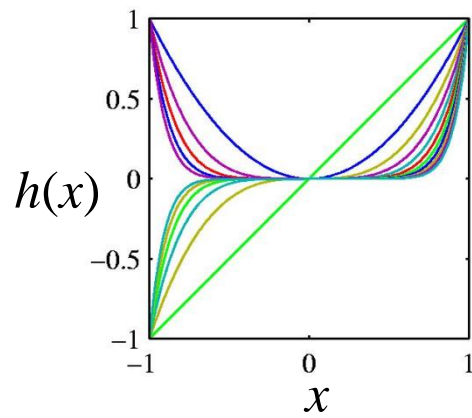
$$f(\vec{x}, \vec{w}) = w_0 + w_1 h_1(\vec{x}) + \cdots + w_m h_m(\vec{x})$$

nonlinear  
relationship  
input  $\longleftrightarrow$  output

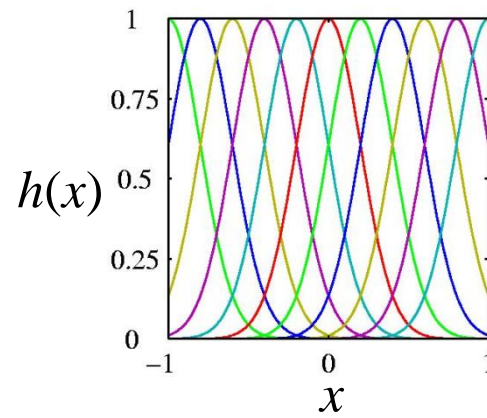
$\{h_i(\vec{x})\}$  is a set of non-linear basis functions (or *kernel* functions)

We can rewrite  $f(\vec{x}, \vec{w}) = \vec{w}^T \begin{bmatrix} 1 \\ h_1(\vec{x}) \\ \vdots \\ h_m(\vec{x}) \end{bmatrix}$

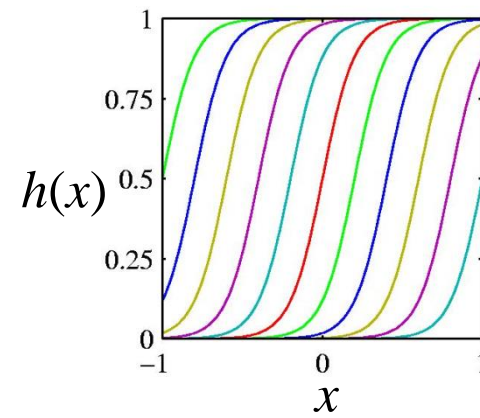
# Examples of Basis Functions



polynomial



Gaussian



sigmoidal

$$\text{Let } \vec{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{bmatrix} \quad X = \begin{bmatrix} 1 & h_1(\vec{x}_1) & h_2(\vec{x}_1) & \cdots & h_m(\vec{x}_1) \\ 1 & h_1(\vec{x}_2) & h_2(\vec{x}_2) & \cdots & h_m(\vec{x}_2) \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & h_1(\vec{x}_N) & h_2(\vec{x}_N) & \cdots & h_m(\vec{x}_N) \end{bmatrix}$$

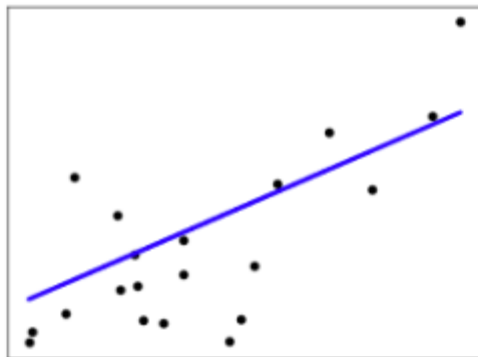
Define the objective function:  $Err = (\vec{y} - X\vec{w})^T (\vec{y} - X\vec{w})$

$$\text{Assume } X \text{ is full rank} \quad \vec{w} = (X^T X)^{-1} X^T \vec{y}$$

# Linear Regression

LinearRegression fits a linear model with coefficients  $w = (w_1, \dots, w_p)$  to minimize the residual sum of squares between the observed targets in the dataset, and the targets predicted by the linear approximation. Mathematically it solves a problem of the form:

$$\min_w ||Xw - y||_2^2$$



`LinearRegression` will take in its `fit` method arrays `x`, `y` and will store the coefficients  $w$  of the linear model in its `coef_` member:

```
>>> from sklearn import linear_model
>>> reg = linear_model.LinearRegression()
>>> reg.fit([[0, 0], [1, 1], [2, 2]], [0, 1, 2])
LinearRegression()
>>> reg.coef_
array([0.5, 0.5])
```





# LinearRegression

```
class sklearn.linear_model.LinearRegression(*, fit_intercept=True, copy_X=True,  
n_jobs=None, positive=False) \[source\]
```

Ordinary least squares Linear Regression.

LinearRegression fits a linear model with coefficients  $w = (w_1, \dots, w_p)$  to minimize the residual sum of squares between the observed targets in the dataset, and the targets predicted by the linear approximation.

## Parameters:

**fit\_intercept** : *bool, default=True*

Whether to calculate the intercept for this model. If set to False, no intercept will be used in calculations (i.e. data is expected to be centered).

**copy\_X** : *bool, default=True*

If True, X will be copied; else, it may be overwritten.

**n\_jobs** : *int, default=None*

The number of jobs to use for the computation. This will only provide speedup in case of sufficiently large problems, that is if firstly `n_targets > 1` and secondly `X` is sparse or if `positive` is set to `True`. `None` means 1 unless in a `joblib.parallel_backend` context. `-1` means using all processors. See [Glossary](#) for more details.

**positive** : *bool, default=False*

When set to `True`, forces the coefficients to be positive. This option is only supported for dense arrays.

#### Attributes:

**coef\_** : *array of shape (n\_features, ) or (n\_targets, n\_features)*

Estimated coefficients for the linear regression problem. If multiple targets are passed during the fit (y 2D), this is a 2D array of shape (n\_targets, n\_features), while if only one target is passed, this is a 1D array of length n\_features.

**rank\_** : *int*

Rank of matrix `X`. Only available when `X` is dense.

**singular\_** : *array of shape (min(X, y),)*

Singular values of `X`. Only available when `X` is dense.

**intercept\_** : *float or array of shape (n\_targets,)*

Independent term in the linear model. Set to 0.0 if `fit_intercept = False`.

**n\_features\_in\_** : *int*

Number of features seen during `fit`.

! Added in version 0.24.

**feature\_names\_in\_** : *ndarray of shape (n\_features\_in\_ ,)*

Names of features seen during `fit`. Defined only when `X` has feature names that are all strings.

# Ridge regression

Ridge regression addresses some of the problems of Ordinary Least Squares by imposing a penalty on the size of the coefficients. The ridge coefficients minimize a penalized residual sum of squares:

$$\min_w ||Xw - y||_2^2 + \alpha ||w||_2^2$$

The complexity parameter  $\alpha \geq 0$  controls the amount of shrinkage: the larger the value of  $\alpha$ , the greater the amount of shrinkage and thus the coefficients become more robust to collinearity.

As with other linear models, `Ridge` will take in its `fit` method arrays `x`, `y` and will store the coefficients  $w$  of the linear model in its `coef_` member:

```
>>> from sklearn import linear_model
>>> reg = linear_model.Ridge(alpha=.5)
>>> reg.fit([[0, 0], [0, 0], [1, 1]], [0, .1, 1])
Ridge(alpha=0.5)
>>> reg.coef_
array([0.34545455, 0.34545455])
>>> reg.intercept_
0.13636...
```

# Ridge

```
class sklearn.linear_model.Ridge(alpha=1.0, *, fit_intercept=True, copy_X=True,  
max_iter=None, tol=0.0001, solver='auto', positive=False, random_state=None) \[source\]
```

Linear least squares with l2 regularization.

Minimizes the objective function:

$$\|y - Xw\|_2^2 + \alpha * \|w\|_2^2$$

This model solves a regression model where the loss function is the linear least squares function and regularization is given by the l2-norm. Also known as Ridge Regression or Tikhonov regularization. This estimator has built-in support for multi-variate regression (i.e., when  $y$  is a 2d-array of shape  $(n\_samples, n\_targets)$ ).

Read more in the [User Guide](#).

## Parameters:

**alpha** : *{float, ndarray of shape (n\_targets,)}, default=1.0*

Constant that multiplies the L2 term, controlling regularization strength. `alpha` must be a non-negative float i.e. in `[0, inf)`.

When `alpha = 0`, the objective is equivalent to ordinary least squares, solved by the `LinearRegression` object. For numerical reasons, using `alpha = 0` with the `Ridge` object is not advised. Instead, you should use the `LinearRegression` object.

If an array is passed, penalties are assumed to be specific to the targets. Hence they must correspond in number.

**fit\_intercept** : *bool, default=True*

Whether to fit the intercept for this model. If set to false, no intercept will be used in calculations (i.e. `X` and `y` are expected to be centered).

**copy\_X** : *bool, default=True*

If True, X will be copied; else, it may be overwritten.

**max\_iter** : *int, default=None*

Maximum number of iterations for conjugate gradient solver. For 'sparse\_cg' and 'lsqr' solvers, the default value is determined by `scipy.sparse.linalg`. For 'sag' solver, the default value is 1000. For 'lbfgs' solver, the default value is 15000.

**tol** : *float, default=1e-4*

The precision of the solution (`coef_`) is determined by `tol` which specifies a different convergence criterion for each solver:

- 'svd': `tol` has no impact.
- 'cholesky': `tol` has no impact.
- 'sparse\_cg': norm of residuals smaller than `tol`.
- 'lsqr': `tol` is set as `atol` and `btol` of `scipy.sparse.linalg.lsqr`, which control the norm of the residual vector in terms of the norms of matrix and coefficients.
- 'sag' and 'saga': relative change of coef smaller than `tol`.
- 'lbfgs': maximum of the absolute (projected) gradient= $\max|\text{residuals}|$  smaller than `tol`.

**solver** : {'auto', 'svd', 'cholesky', 'lsqr', 'sparse\_cg', 'sag', 'saga', 'lbfgs'}, default='auto'

Solver to use in the computational routines:

- 'auto' chooses the solver automatically based on the type of data.
- 'svd' uses a Singular Value Decomposition of X to compute the Ridge coefficients. It is the most stable solver, in particular more stable for singular matrices than 'cholesky' at the cost of being slower.
- 'cholesky' uses the standard `scipy.linalg.solve` function to obtain a closed-form solution.
- 'sparse\_cg' uses the conjugate gradient solver as found in `scipy.sparse.linalg.cg`. As an iterative algorithm, this solver is more appropriate than 'cholesky' for large-scale data (possibility to set `tol` and `max_iter`).
- 'lsqr' uses the dedicated regularized least-squares routine `scipy.sparse.linalg.lsqr`. It is the fastest and uses an iterative procedure.
- 'sag' uses a Stochastic Average Gradient descent, and 'saga' uses its improved, unbiased version named SAGA. Both methods also use an iterative procedure, and are often faster than other solvers when both `n_samples` and `n_features` are large. Note that 'sag' and 'saga' fast convergence is only guaranteed on features with approximately the same scale. You can preprocess the data with a scaler from `sklearn.preprocessing`.
- 'lbfgs' uses L-BFGS-B algorithm implemented in `scipy.optimize.minimize`. It can be used only when `positive` is True.

All solvers except 'svd' support both dense and sparse data. However, only 'lsqr', 'sag', 'sparse\_cg', and 'lbfgs' support sparse input when `fit_intercept` is True.

**positive** : bool, default=False

When set to `True`, forces the coefficients to be positive. Only 'lbfgs' solver is supported in this case.

**random\_state** : int, RandomState instance, default=None

Used when `solver` == 'sag' or 'saga' to shuffle the data. See [Glossary](#) for details.

#### Attributes:

**coef\_** : ndarray of shape (n\_features,) or (n\_targets, n\_features)

Weight vector(s).

**intercept\_** : float or ndarray of shape (n\_targets,)

Independent term in decision function. Set to 0.0 if `fit_intercept = False`.

**n\_iter\_** : None or ndarray of shape (n\_targets,)

Actual number of iterations for each target. Available only for sag and lsqr solvers. Other solvers will return None.

! Added in version 0.17.

**n\_features\_in\_** : int

Number of features seen during `fit`.

! Added in version 0.24.

**feature\_names\_in\_** : ndarray of shape (n\_features\_in\_,)

Names of features seen during `fit`. Defined only when `X` has feature names that are all strings.

! Added in version 1.0.

**solver\_** : str

The solver that was used at fit time by the computational routines.

! Added in version 1.5.



