






Chapter 10 – NN

Instructor: Dr. Hongfu Liu
Email: hongfuliu@brandeis.edu




Mathematical Model of a Neuron



$$y = f(\vec{x}) = \phi \left(\sum_{i=1}^p w_i x_i + b \right)$$


Mathematical Model of a Neuron




$$y = f(\vec{x}) = \phi(\vec{w}^T \vec{x})$$

Activation Functions


$$f(\vec{x}) = \phi(\vec{w}^T \vec{x})$$



Sigmoid



tanh




relu

Activation Functions

Sigmoid Function


$$\phi(v) = \frac{1}{1 + \exp(-av)} \quad a > 0, -\infty < v < \infty$$



Activation Functions

Hyperbolic Tangent Function


$$\phi(v) = a \tanh(bv) = a \frac{e^{bv} - e^{-bv}}{e^{bv} + e^{-bv}} \quad a > 0, b > 0$$




Activation Functions

ReLU (rectified linear unit)

$$\phi(x) = x^+ = \max(0, x) = \begin{cases} x & \text{if } x > 0 \\ 0 & \text{otherwise} \end{cases}$$



Error-Correction Learning



If the error function is continuously differentiable, we can use gradient descent


$$\Delta \vec{w} = \eta \frac{\partial e_{\vec{w}}(y, \phi(v))}{\partial \vec{w}}$$

If the activation function ϕ is differentiable, the necessary condition for minimize the error is

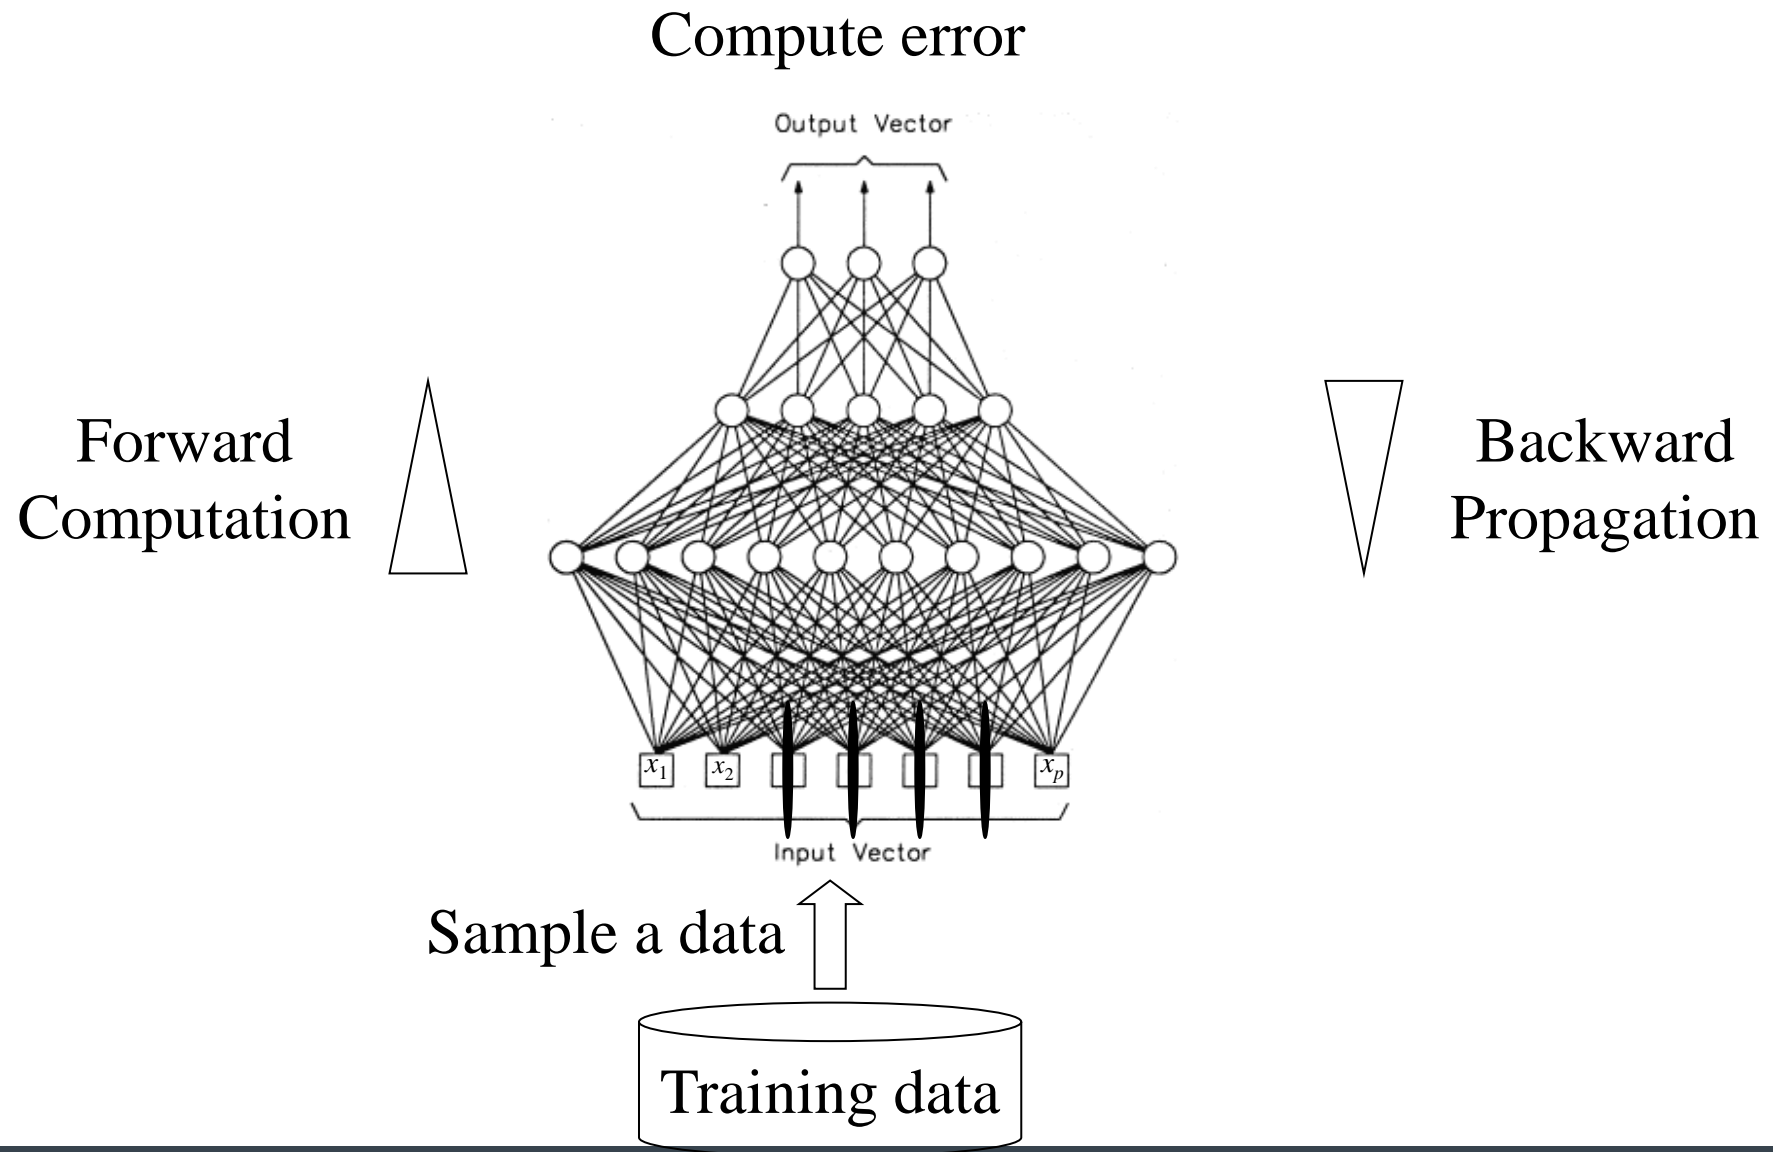
$$\frac{\partial e_{\vec{w}}(y, \phi(v))}{\partial \vec{w}} = \begin{bmatrix} \frac{\partial e_{\vec{w}}(y, \phi(v))}{\partial w_0} \\ \frac{\partial e_{\vec{w}}(y, \phi(v))}{\partial w_1} \\ \vdots \\ \frac{\partial e_{\vec{w}}(y, \phi(v))}{\partial w_p} \end{bmatrix} = 0$$

Hence, we can use gradient-based search algorithm to train a neuron.

Multilayer Neural Networks



The Back-Propagation Algorithm



Sequential vs Batch

- Sequential mode
- Batch mode

$$Err = \frac{1}{2N} \sum_{n \in \text{samples}} \sum_{\substack{j \in \text{output} \\ \text{neurons}}} e_{nj}^2$$

$$\Delta w_{ji} = -\eta \frac{\partial Err(t)}{\partial w_{ji}(t)} = \eta \sum_n e_{nj} \frac{\partial e_{nj}}{\partial w_{ji}}$$

Online, memory, implementation, stochastic

- Number of hidden neurons
- Sequential training mode
- Batch training mode

Convolution (Conv)


- Local receptive field



1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

1	0	1
0	1	0
1	0	1

- Spatial invariant




1x1	1x0	1x1	0	0
0x0	1x1	1x0	1	0
0x1	0x0	1x1	1	1
0	0	1	1	0
0	1	1	0	0

4		

- Convolution can be efficiently computed


Padding and Stride



Input size is 4×4 , conv filter size is 3×3 , padding = 2
What is the output size?

Conv Layers


A Conv layer contains multiple convolutional neurons (or filters, feature extractors, ...)



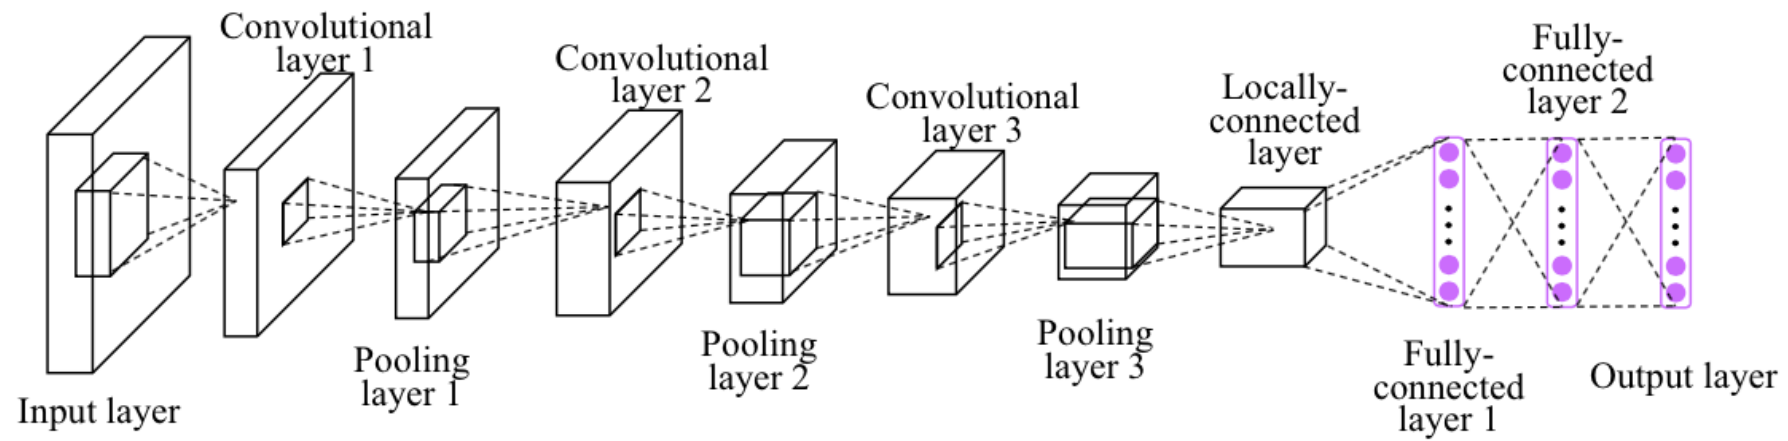
The feature maps produced by a convolution layer of 10 neurons

By one of the neurons By another neuron


Pooling Layer




A Deep CNN Example



Residual Neural Network (ResNet)

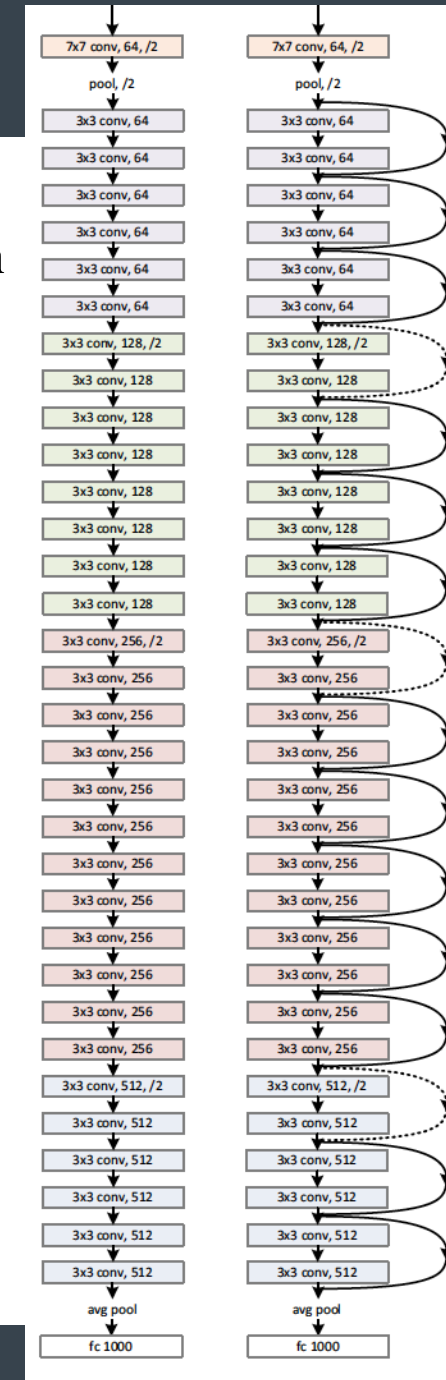
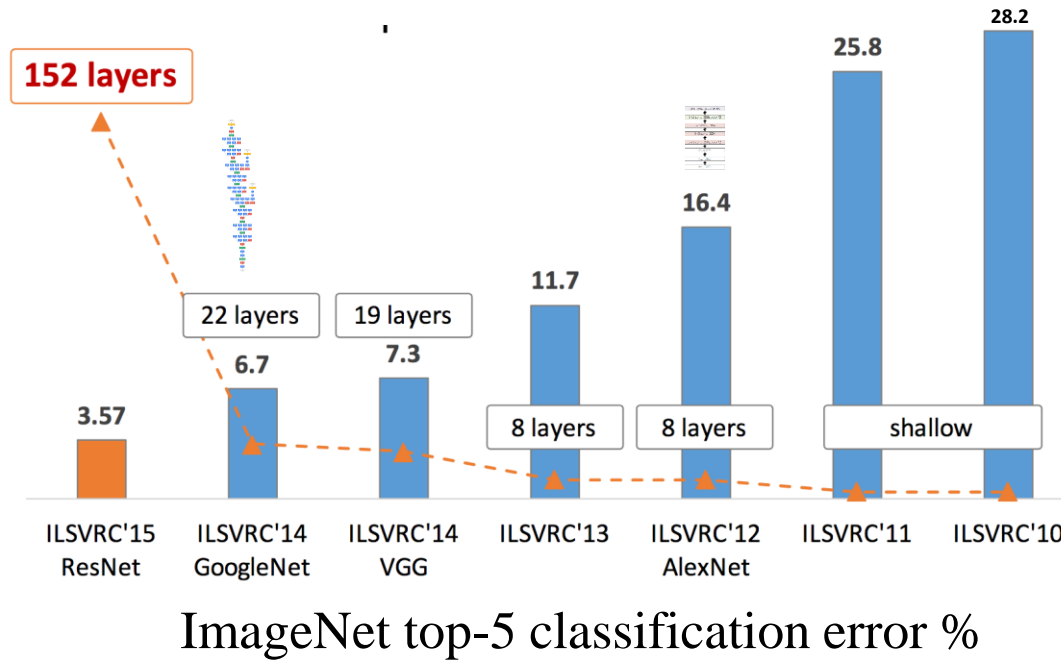


Source: wikipedia




Residual Block

ResNets – Go Deeper




Densely Connected Convolutional Networks

(DenseNet)

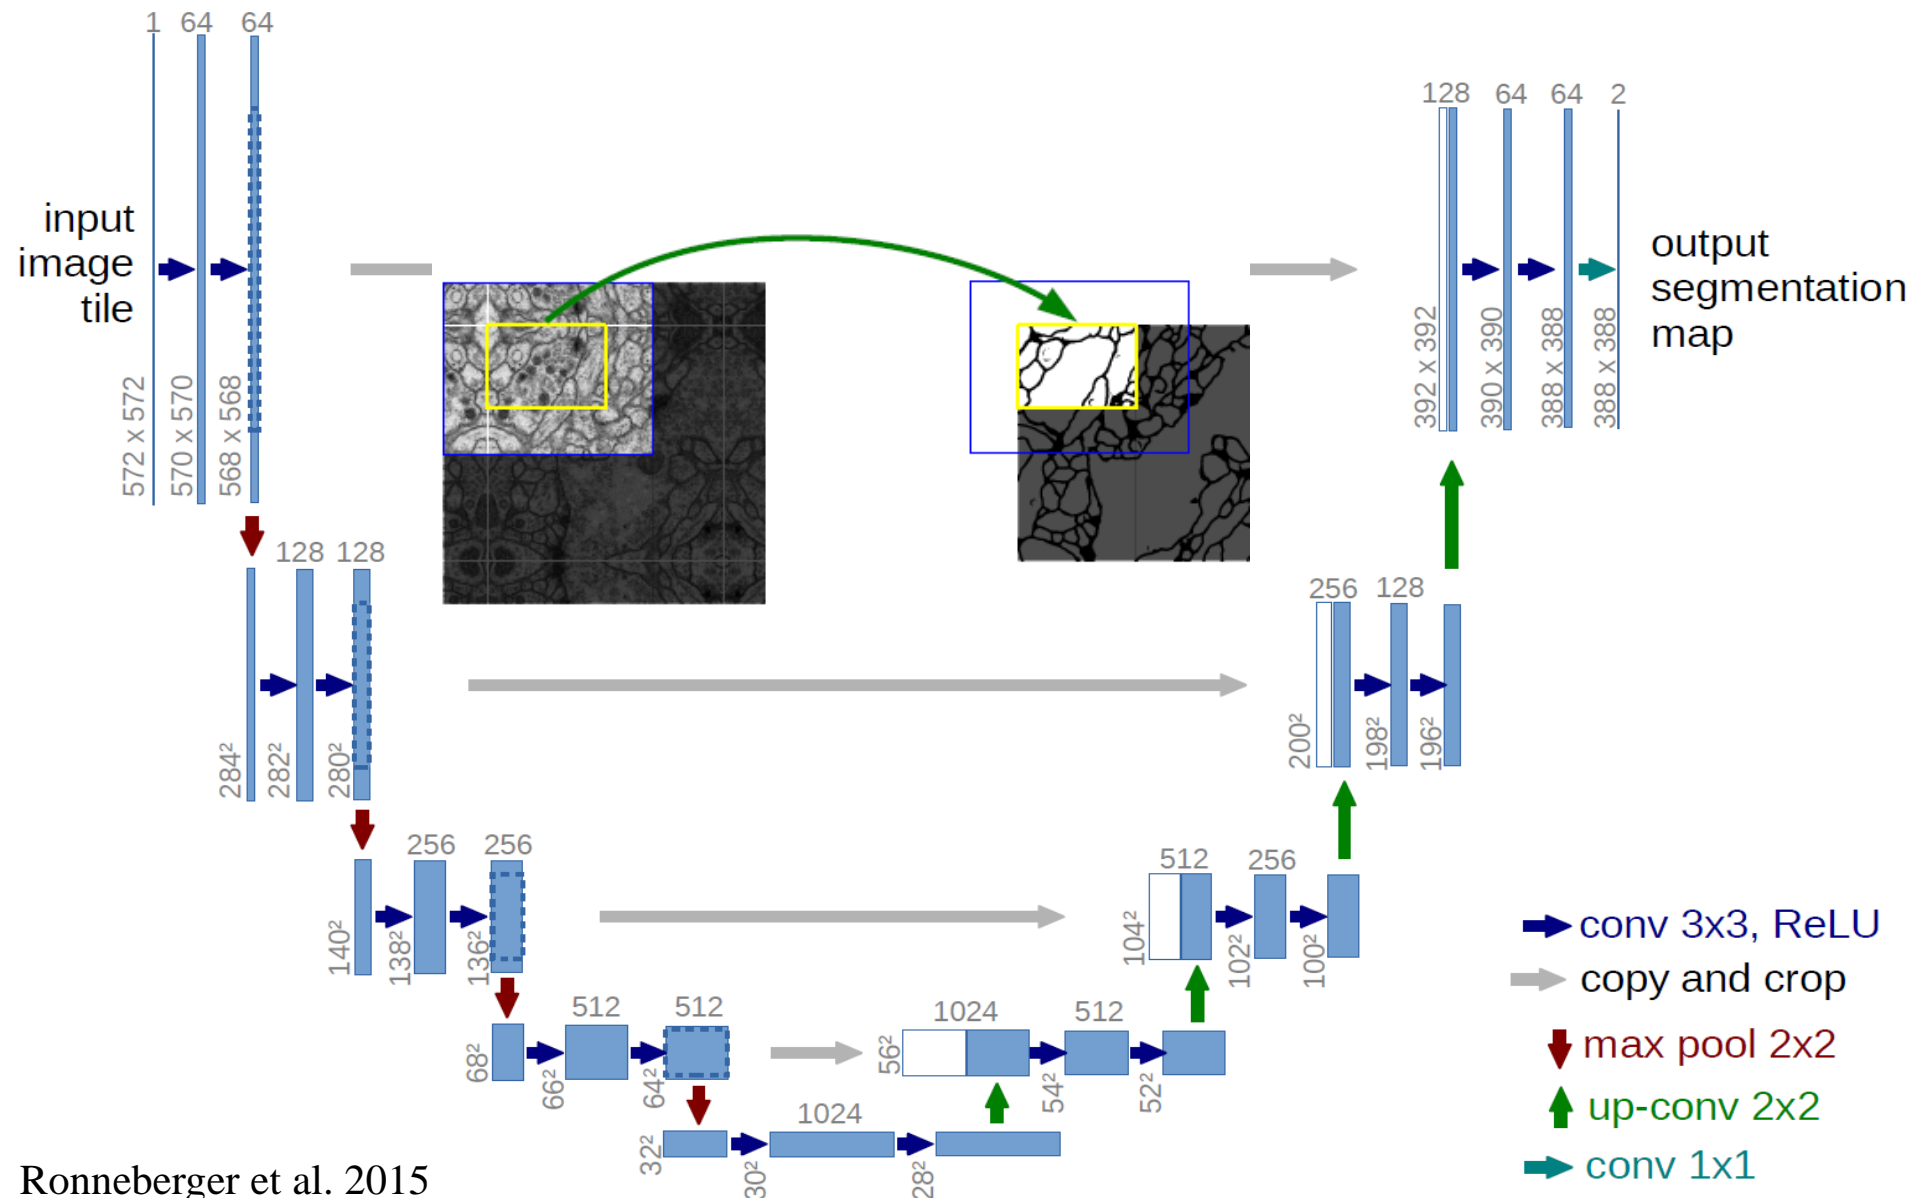


Residual block

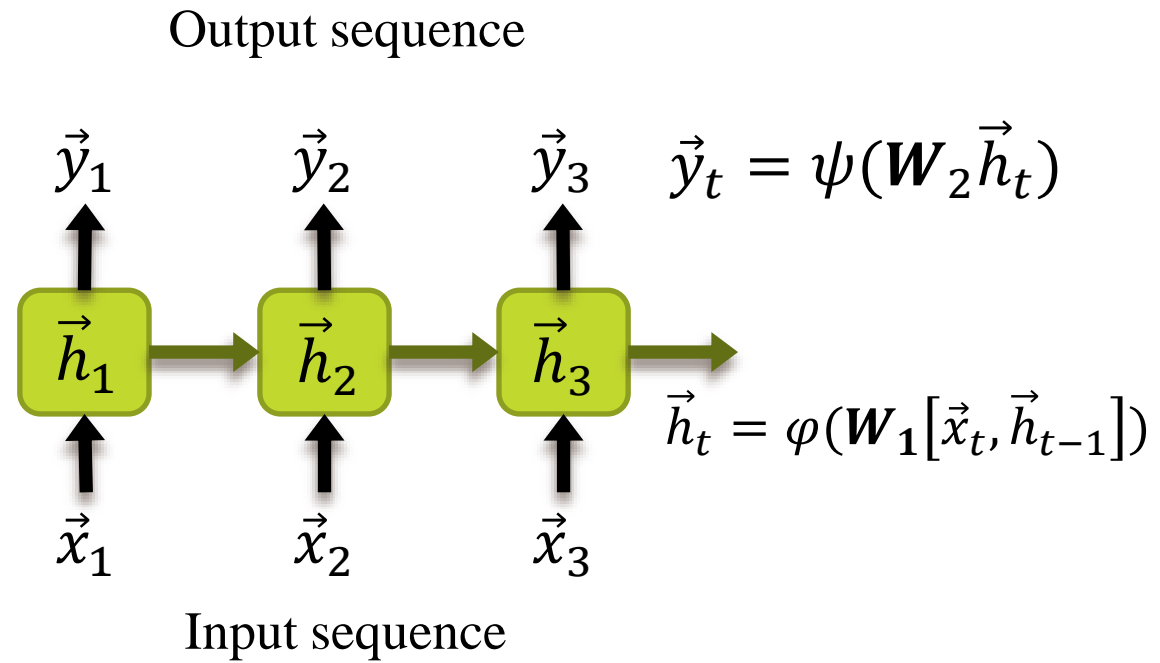


Dense block

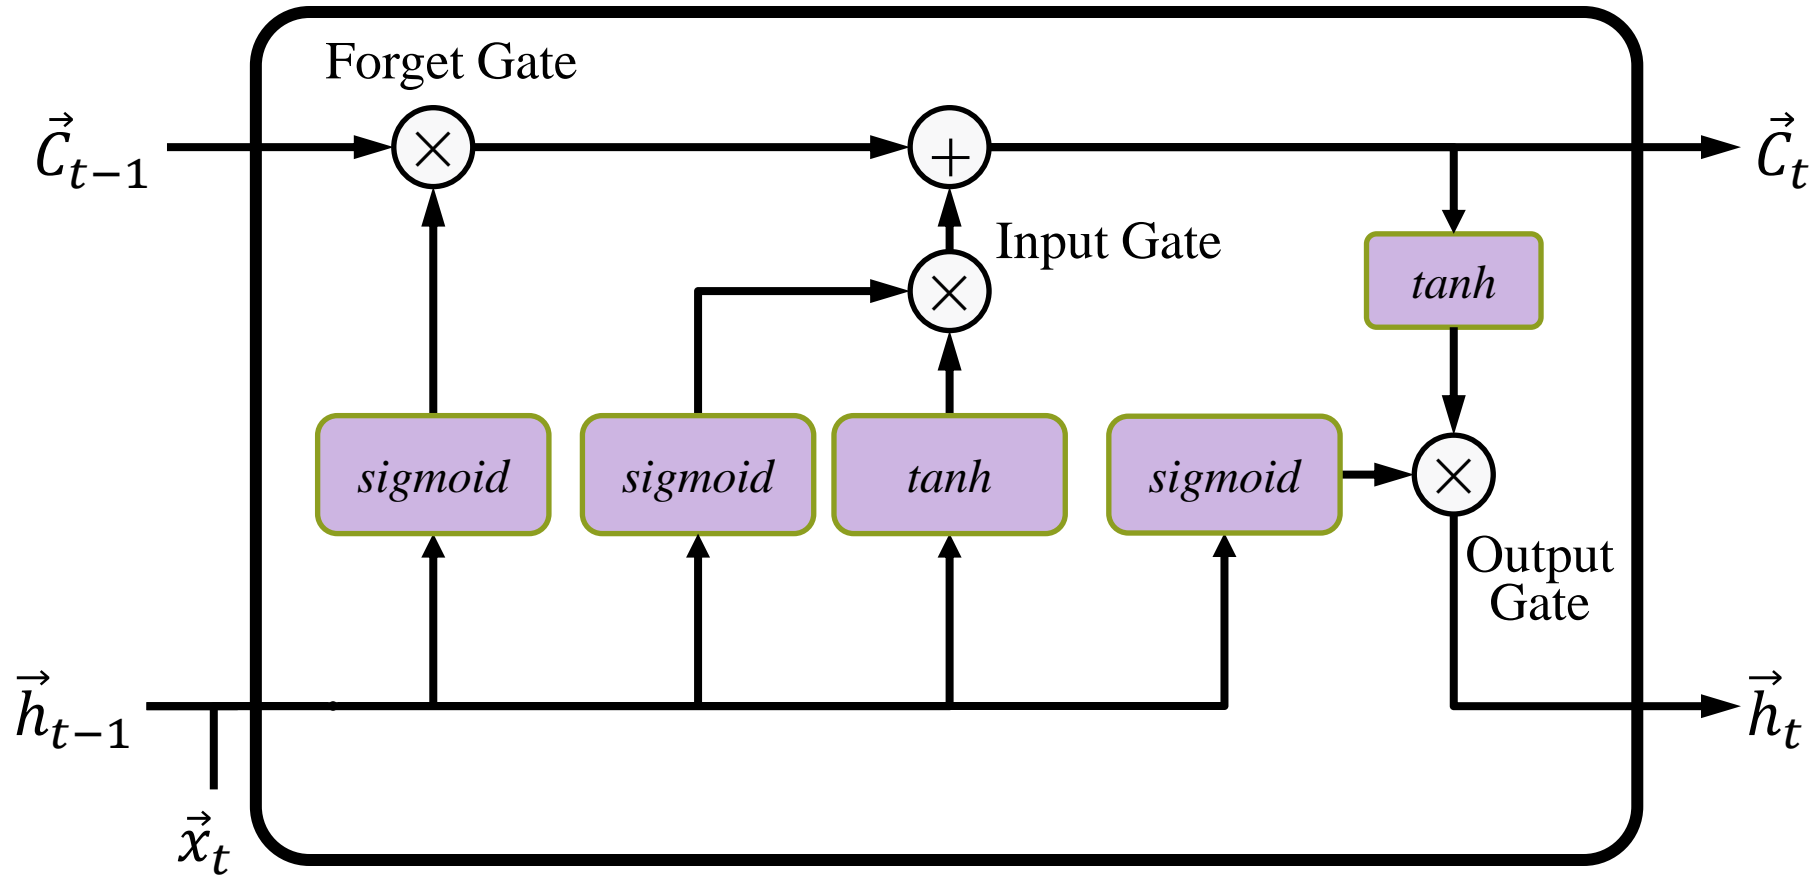
U-Net (for image segmentation)



Recurrent Neural Networks



Long Short-Term Memory (LSTM) Cell



Gated Recurrent Unit

